

Bullet-Fluids Reference (draft version)

02 January 2014

Abstract

Originally based on Fluids v.2 [2], the Bullet-Fluids library is a Zlib licensed extension to the Bullet Physics engine [1] that is targeted at interactive or real-time simulation of fluids using the meshless particle method known as SPH.

The library is being developed as a draft for the production version, which is targeted at Bullet 3.x.

This document provides an introduction to the Bullet-Fluids library, and also includes some explanation of parameters and other details.

Contents

1	Introduction to Bullet-Fluids	3
1.1	Important notes (please read this section before using the library)	3
1.1.1	Simulation Scale	3
1.1.2	Changing Particle Indices	3
1.1.3	Explosions	4
1.2	Class Hierarchy	4
1.2.1	btFluidRigidDynamicsWorld	4
1.2.2	btFluidSph	4
2	A SPH Fluid Primer	5
2.1	A simple particle simulation	5
2.2	Extending the simple particle simulation to fluids	7
3	Parameters	11
3.1	Global Parameters	11
3.2	Local Parameters	11
3.3	Other Parameters	13
4	Optimizations	14
4.1	Grid Algorithm	14
4.1.1	Static Noncolliding Uniform Grid (implemented by Bullet-Fluids)	14
4.2	Grid Cell Size	15
4.3	GPU Optimizations	15
4.3.1	3-cell bar	15
4.3.2	Effect of CPU optimizations on the GPU	15
4.4	CPU Optimizations	16
4.4.1	Symmetry and Multithreading	16
4.4.2	Neighbor Tables	17

5	Other Internal Details	18
5.1	Update Loop	18
5.2	Rigid Body Interaction	18
5.3	Other notes	18

1 Introduction to Bullet-Fluids

1.1 Important notes (please read this section before using the library)

1.1.1 Simulation Scale

The fluid simulation defines two scales, *world scale* and *fluid simulation scale* (hereafter, *simulation scale*).

- *simulation scale* is the scale at which the SPH density and forces are calculated.
- *world scale* is the scale at which rendering, the rigid body simulation and all other functions take place.

In general, the world scale is much larger than the simulation scale. Multiplying a value by the simulation scale can be seen as scaling down the value from world or rigid body units to fluid units, while dividing a value by the simulation scale is equivalent to scaling up the value from the fluid simulation scale to world units.

$$\text{simulation_scale.length} = \text{world_scale.length} * \text{simulation_scale}$$

$$\text{world_scale.length} = \text{simulation_scale.length} / \text{simulation_scale}$$

To give a concrete example, the default simulation scale is 0.004.

This means that a length of 1 metre at world scale is shrunk to 0.004 metres at simulation scale.

Alternatively, a length of 1 metre at simulation scale is expanded to 250 metres at world scale.

In either case, the ratio of world scale to simulation scale is (1 / simulation_scale), or 1 : 250 by default.

When changing the size of the particles, it is recommended to change `m_simulationScale` and parameters marked [world scale]. Reducing the simulation scale will make particles larger, and increasing it will make the particles smaller. For example, doubling the default simulation scale of 0.004 will halve the size of the particles. (default: $1 / 0.004 = 250$; halved: $1 / 0.008 = 125$).

Although in theory it should only be necessary to change the simulation scale to resize the fluid, in practice floating point rounding can cause the fluid to behave differently when the simulation scale is changed. (Even if there is no difference mathematically.) However, this should not matter unless the simulation scale is changed by a factor of 100-1000x or greater.

The motivation for including the simulation scale parameter is to provide a easy way to change the scale of the fluid without affecting its behavior. The Navier-Stokes equations are sensitive to scale, and so the SPH approximation of their terms is also scale sensitive.

Additionally, SPH is highly sensitive to the parameters used; if a single parameter is even slightly incorrect, it can cause the entire simulation to explode. Abstracting the size of the fluid using a ‘simulation scale’ term allows us to shrink or expand the fluid without spending excessive effort on tuning the parameters. Were the simulation scale not implemented, it would be a difficult task to get all parameters correct on the first try. Using two scales allows us to make slight modifications from a working set of parameters to get the desired behavior.

Some more details on the simulation scale are provided at the Fluids v.2 website [2].

1.1.2 Changing Particle Indices

The particle data is resorted every frame. For instance, the particle at index 0 could in the next frame be at index 3, or some other location. This means that the index of each particle cannot be used to track it. However, each particle has a user pointer that can be used to associate a unique id or struct with

it. The per-particle user pointer can be accessed through `btFluidSph::getParticleUserPointer()` and `btFluidSph::setParticleUserPointer()`. This user pointer is set to 0 when the particle is created, and is sorted along with the particle data every frame. It is not modified otherwise.

The main reason for rearranging the particle data is performance; sorting the particles increases the cache hit rate.

1.1.3 Explosions

A common issue that may be encountered when implementing a SPH fluid simulation is the ‘explosion’, where the particles cease to behave as a fluid and begin to fly around everywhere.

This issue is usually caused by applying very high forces to the fluid particles. The source of such forces is many; it could occur if any of the following is set too high: time step, stiffness, gravity, viscosity or other user applied forces.

1.2 Class Hierarchy

There are 2 main classes, and those 2 classes contain 5 important subclasses.

1.2.1 `btFluidRigidDynamicsWorld`

Contains all objects (rigid bodies, collision objects, and fluids) in the simulation.

- `btFluidSphParametersGlobal` - Properties shared by all fluids.
- `btFluidSphSolver` - Determines how the fluid particles interact with each other, and contains solver-specific data.

1.2.2 `btFluidSph`

A single `btFluidSph` corresponds to a group of particles with the same characteristics.

- `btFluidParticles` - Contains particle state data: position, velocity, accumulated force, etc.
- `btFluidSphParametersLocal` - Defines the fluid material type; contains the fluid’s viscosity, particle mass, and other parameters related to SPH and rigid body interaction.
- `btFluidSortingGrid` - A uniform grid used to accelerate the SPH density and force calculation.

2 A SPH Fluid Primer

This covers some of the theory behind SPH, see [Introduction to Bullet-Fluids](#) for library-specific details.

Smoothed Particle Hydrodynamics (SPH) is a method that can be used to simulate fluids (liquids and gases) using particles. Before explaining SPH, however, it is necessary to go over some basic details of particle systems.

Prerequisites and notation: Familiarity with differential calculus and time stepping is very useful, but not absolutely necessary to understand the main points. Nonbold letters are scalars, while bold letters are vectors. For example, a 3-dimensional point using the Cartesian Coordinate System is written as: $\mathbf{r} = (r_x, r_y, r_z)$. ∇ denotes the gradient. ∇^2 is used to represent the Laplacian or the divergence of the gradient; it also appears as $\nabla \cdot \nabla$. The partial derivative of a variable with respect to time is written as $\frac{\partial}{\partial t}$.

2.1 A simple particle simulation

Consider a simple simulation of rigid body spheres without rotation. The rigid spheres are represented as particles with the properties: mass, radius, position, and velocity.

To perform the simulation, we will use a technique known as *time stepping*. We discretize time into finite amounts, called *frames*, and also store some information about each particle, the *state*. The state includes values that change over time; in particular, the position and velocity of each particle. To begin, we initialize each particle with a state. Every frame, we advance the simulation by a fixed amount of time, the *time step*, and update each particle's position and velocity. In order to prevent the spheres from intersecting, we modify the velocities of the particles by applying forces so that they will not penetrate in the next frame.




Figure 1: Example of time stepping for a falling sphere. Time is split into discrete chunks, called *frames*, and each frame advances the simulation by a *time step* of 16 milliseconds (ms).

The update loop for such a particle system might be written as:

- Detect collisions
- Determine collision response forces
- Integrate velocity(that is, apply collision response forces and other forces such as gravity)
- Integrate position

Let's go over that loop in a bit more detail:

Detect collisions The main purpose of this stage is to generate information needed for the collision response stage. For each pair of particles, we calculate a normal vector and distance. The normal vector, which is a vector of length 1 pointing from one particle to another, determines the direction of the repulsive force used to separate the particles. The distance determines whether a force needs to be applied and, if below 0 (which means that the particles are penetrating), affects the magnitude of the force that is applied.

Determine collision response forces In this stage, we generate forces that will change the velocities of the particles so that, ideally, they will not interpenetrate in the next step. The force is called the *normal force* (as it is applied along the normal vector), which scales up as large as needed to prevent the particles from penetrating. It is also possible to generate a friction force, so that the particles will not slide across each other.




Figure 2: The collision detection stage generates distance and a normal vector for each pair of colliding spheres. If the spheres are able to rotate, then it is also necessary to detect the point of contact where the spheres are touching. Whether the normal vector points towards the left or right sphere is a matter of convention.

Integrate velocity and position In these stages, we advance the state of the simulation, updating the velocity and position of each particle.

$$\text{velocity}_{\text{next}} = \text{velocity} + (\text{force}/\text{mass}) * \text{timestep}$$

$$\text{position}_{\text{next}} = \text{position} + \text{velocity}_{\text{next}} * \text{timestep}$$

It is important to note that we use $\text{velocity}_{\text{next}}$ to compute the next position as opposed to the current velocity . This means that the velocity and position are integrated using semi-implicit Euler as opposed

to explicit Euler. The difference is subtle but very important, as explicit Euler is unconditionally unstable; that is, it ensures that the simulation will explode if damping (artificially reducing the velocity) is not applied.

2.2 Extending the simple particle simulation to fluids

Smoothed Particle Hydrodynamics may sound intimidating, but it is actually not very different from this particle simulation. In particular, the only difference between a SPH fluid simulation and a rigid body particle simulation is that the *determine collision response forces* stage is different.

In particular, the rigid body simulation uses a repulsive force between colliding particles, while an SPH fluid uses pressure and viscosity forces based on the Navier-Stokes equations. This has the effect of changing some characteristics of the particles. Instead of considering the particles as solid spheres of uniform density, with the mass evenly distributed across the particle, the particles are seen as points with a radius of interaction. Furthermore, the boundary between SPH particles is soft. A common rule of thumb for SPH simulations is that each particle should have 20–40 neighbors; that is, each particle is expected to be interpenetrating with 20 to 40 other particles. Finally, the mass of a SPH particle is unevenly distributed throughout the volume of the particle, with more mass concentrated at the center and less mass farther out.

Another difference is that while rigid body particles are seen as distinct, separate entities, SPH particles are viewed as being part of a fluid, with each SPH particle contributing to a greater whole. In particular, each SPH particle contributes some mass to a density scalar field. That is, it is possible to use particles to associate every point in space with a density. Just as it is possible to use linear interpolation to define a line from 2 points, it is possible to use SPH to define a scalar field from many points.

To reiterate, in addition to the properties of mass, radius, position, and velocity, SPH particles also gain the implicit property of density. This property is not stored with each particle, but interpolated by summing the contributions from all particles.




Figure 3: On the left are uniform density rigid particles with hard boundaries. On the right are distributed density SPH particles with soft boundaries. The darkness indicates the amount of mass at that point.

The particular equation we use to define the density, ρ , at a position, \mathbf{r} , is:

$$\rho(\mathbf{r}) = \sum_j m_j W_{poly6}(\mathbf{r} - \mathbf{r}_j, h) \quad (1)$$

Where the sum loops over all nearby particles with position \mathbf{r}_j within the radius of interaction, h , the term m_j is the mass of a particle. W is used to denote a *smoothing kernel*; it is a function that determines how a property is distributed throughout the volume of a fluid particle. In this case, we are using the kernel named the *poly6* kernel. In the context of our simulation, W_{poly6} is a function that determines how the mass is distributed throughout a fluid particle. The amount of mass contributed is highest at the particle's center, and gradually falls to 0 as the distance from the center increases to the radius of interaction, h .




Figure 4: Example of equation 1 in 1D. SPH Interpolates a global density(blue) by summing up local densities(black) from particles.

Fluid motion is described by the the Navier-Stokes equations. Although the Navier-Stokes equations are a set of equations, we are only interested in a single equation due to various assumptions and other simplifications:

$$\rho \frac{\partial \mathbf{v}}{\partial t} = -\nabla p + \mu \nabla^2 \mathbf{v} + \mathbf{f} \quad (2)$$

With density ρ , acceleration $\partial \mathbf{v} / \partial t$, pressure p , and velocity \mathbf{v} .

This equation may appear complex, but it is actually quite simple; it is just $F = ma$ for fluids.

- $-\nabla p$, the negated gradient of pressure, means that fluids are accelerated from places of higher pressure towards areas of lower pressure,
- $\mu \nabla^2 \mathbf{v}$, the Laplacian of velocity, means that fluid particles experience friction against each other, with μ being a constant that determines the amount of friction, and
- \mathbf{f} includes external forces such as gravity, surface tension and collision response forces (for instance, forces from collisions with rigid bodies).

The main term in this equation that accounts for the motion of fluids such as water is the pressure term, $-\nabla p$. While the normal force that is applied between rigid body particles acts as a hard boundary that removes all penetration, the pressure force acts to eliminate variations in pressure. As we shall see next, eliminating variations in pressure also has the effect of constraining the particles' density to a single value.

Minor note: although the Laplacian operates on a scalar field, applying it to a vector simply means to apply it to each of the components. $\nabla^2 \mathbf{v} = (\nabla^2 v_x, \nabla^2 v_y, \nabla^2 v_z)$, where $\mathbf{v} = (v_x, v_y, v_z)$.

The density scalar field can be converted into a pressure field by using an equation of state. For liquids, the equation we use is:

$$p = k(\rho - \rho_0) \quad (3)$$

Where k , the stiffness, is a positive constant that determines the strength of the pressure force and ρ_0 is the fluid's rest density. The *rest density* is exactly what it sounds like—the density of the fluid when it is at *rest*; that is, when it has settled and is not moving.

The effect of this equation is that particles below ρ_0 , the rest density, will have a negative pressure, while particles above the rest density will have a positive pressure. As the gradient operator points in the direction of greatest increase, this results in a force that pushes the particles apart when above rest density, and together when below rest density.

Now that we have a way of defining a pressure field, we have all the variables needed to approximate a solution for the acceleration using SPH.

$$acceleration = \frac{\partial \mathbf{v}}{\partial t} = \frac{-\nabla p + \mu \nabla^2 \mathbf{v} + \mathbf{f}}{\rho} \quad (4)$$

Using SPH, the term $-\nabla p$ is approximated at a particle i as:

$$\mathbf{f}_i^{pressure} = - \sum_j m_j \frac{p_i + p_j}{2\rho_j} \nabla W_{spiky}(\mathbf{r}_i - \mathbf{r}_j, h) \quad (5)$$

and the viscosity term $\mu \nabla^2 \mathbf{v}$ as:

$$\mathbf{f}_i^{viscosity} = \mu \sum_j m_j \frac{\mathbf{v}_j - \mathbf{v}_i}{\rho_j} \nabla^2 W_{viscosity}(\mathbf{r}_i - \mathbf{r}_j, h) \quad (6)$$

Where ∇W_{spiky} is the gradient of the spiky kernel, and $\nabla^2 W_{viscosity}$ is the Laplacian of the viscosity kernel. The spiky and viscosity kernels are functions that have been specially designed to improve stability when using SPH to approximate pressure and viscosity forces. Similar to the W_{poly6} kernel, they scale the force between pairs of particles, with the force gradually dropping to 0 as the distance between particles increases to the radius of interaction.

Substituting terms from the kernels W_{poly6} , ∇W_{spiky} , and $\nabla^2 W_{viscosity}$, removing constant terms from the sum, and then solving for acceleration ($\mathbf{f}_i = \rho_i \mathbf{a}_i$) changes the equations into this form:

$$\rho_i(\mathbf{r}_i) = \frac{315m}{64\pi h^9} \sum_j (h^2 - |\mathbf{r}_i - \mathbf{r}_j|)^3 \quad (7)$$

$$\mathbf{a}_i^{pressure} = \frac{45m}{\pi h^6} \sum_j \frac{p_i + p_j}{2\rho_i \rho_j} \frac{\mathbf{r}_i - \mathbf{r}_j}{|\mathbf{r}_i - \mathbf{r}_j|} (h - |\mathbf{r}_i - \mathbf{r}_j|)^2 \quad (8)$$

$$\mathbf{a}_i^{viscosity} = \frac{45\mu m}{\pi h^6} \sum_j \frac{\mathbf{v}_j - \mathbf{v}_i}{\rho_i \rho_j} (h - |\mathbf{r}_i - \mathbf{r}_j|) \quad (9)$$

Going over the terms again, m is the mass of a particle (assuming all particles have the same mass), h is the radius of interaction, p is pressure, ρ is density, μ is the strength of viscosity, \mathbf{r} is position, and \mathbf{v} is velocity. The subscripts i and j are used to indicate the indices of the particles that the properties are associated with. $|\mathbf{r}_i - \mathbf{r}_j|$ denotes the distance between points \mathbf{r}_i and \mathbf{r}_j .

Every frame, we iterate over all particles twice. The first pass to calculate the density and pressure (equations 7 and 3), and the second pass to calculate pressure and viscosity forces (equations 8 and 9). Finally, we apply these forces to each particle, and integrate their positions.

To conclude, there is only a minor conceptual difference between a rigid body particle simulation and a SPH fluid simulation. In order to convert a rigid body particle simulation into a SPH fluid simulation, we need only to replace the repulsive force with pressure and viscosity forces.

The important things to remember are that:

- Fluid motion is described by the Navier-Stokes equations, which define a pressure and viscosity force.
- SPH is a method that combines the local contributions of particles to represent a global scalar field, and can be used to approximate solutions to the Navier-Stokes equations.
- The SPH pressure force is both attractive and repulsive, and accelerates particles towards the fluid's rest density.
- The SPH viscosity force accelerates particles towards having the same velocity. This has a similar effect to friction, dissipating energy and gradually reducing the velocity of the particles. (2 particles with the same velocity experience no acceleration from the viscosity force, while particles moving in opposite directions experience a strong viscosity force.)

The derivation presented in this article was originally introduced by Müller, Charypar, and Gross in [4]. It is recommended to see their paper for a more in depth explanation. Additionally, a good tutorial on SPH is provided in Kelager's Master's thesis [3].

3 Parameters

In the context of real time simulations, factors such as the time step and low stiffness mean that the relation to physical units becomes largely academic, so using actual physical values is unlikely to give the desired result. Parameters at world scale are marked as [world scale] and parameters at simulation scale are marked as [simulation scale]. [X, Y] means that the range of the parameter should be between X and Y. For instance, [0.0, 1.0] denotes a range from 0 to 1.

3.1 Global Parameters

Properties assigned to all `btFluidSph` in a `btFluidRigidDynamicsWorld`.

m.timeStep For performance and stability reasons, it is not possible to run the fluid simulation at the same time step as the rigid body simulation. Setting this value too high will cause the fluid to explode. In general, a higher time step requires lower stiffness to prevent the fluid from exploding. As a rule of thumb, doubling the time step means that the stiffness needs to be halved. By default, the fluid is stepped 5 times slower(3ms) than the rigid body simulation(16ms).

m.simulationScale See [1.1.1](#).

m.gravity [simulation scale] Acceleration applied to all particles during each frame; by default the Y-axis is assumed to point downwards. Note that, as this is a simulation scale parameter, the default gravity is 250x larger compared to rigid bodies. This default value can be reduced considerably. (Part of the reason for the high default gravity is the ratio of time steps between fluid/rigid).

m.sphSmoothRadius [simulation scale] Commonly denoted as r or h in the literature. Defines the radius in which SPH particles will interact with each other. This should be set so that each particle has about 20–40 neighbors. To be precise, a particle influences another if the distance to its center is less than or equal to r —that is, it should be thought of as sphere v. point collision. Considering it as sphere-sphere collision means that both spheres would have a radius of $r/2$.

3.2 Local Parameters

Properties assigned to a single `btFluidSph`.

m.viscosity Affects the magnitude of the viscosity force, which accelerates particles so that they have the same velocity—it defines the resistance of the fluid to flow. For instance, water has low viscosity while honey has high viscosity. The viscosity force also works to improve the numerical stability of the simulation by reducing particle velocities. However, a too high viscosity value will instead introduce energy into the simulation.

m.restDensity Defines the density of the fluid when it is at rest, or not moving. Water has a rest density of $1000\text{kg}/\text{m}^3$. The pressure force is applied to constrain particles to this density.

m.sphParticleMass Contribution of a single particle to the SPH density scalar field.

m_stiffness [simulation scale] The fluid stiffness determines how violently the fluid reacts to variations in density. In other words, it defines the magnitude of the SPH pressure force. For mostly incompressible fluids such as water, is best to keep this value as high as possible. Another reason to use a higher stiffness value is performance; a particle that is part of a compressible (low stiffness) fluid tends to have more neighbor particles. The maximum stiffness that can be used is dependent on the time step.

m_initialSum (0.0, 1.0) Determines the particle's self contribution when calculating its density. It should always be greater than 0, but generally be kept as low as possible. Using a value of 1 is equivalent to adding a virtual SPH particle at zero distance when calculating each particle's density; the virtual particle is not included when calculating forces.

Higher values improve stability, especially for particles that are isolated or have incomplete neighborhoods. As it has the effect of increasing the density of all particles, however, it also tends to make particles repulse each other more strongly.

m_particleDist [simulation scale] This is an automatically calculated value that defines a default spacing for particles. It has no effect on the simulation, but should not be changed. Consider spawning a group of fluid particles in the shape of a box. Placing the fluid particles too far apart will cause them to attract (box implodes), while spacing them too closely will make them repulse (box explodes). Depending on how m_initialSum is set, this may need to be scaled by a factor of 0.75–1.25.

m_particleRadius [world scale] Radius of a particle with colliding with btRigidBody/btCollisionObjects. The particle is treated as a rigid body sphere. Has no effect on the SPH simulation.

m_particleMargin [0.0, m_particleRadius) [world scale] Meters of allowed penetration when colliding with rigid bodies. This value should be a small fraction of m_particleRadius. Should not be set to 0, as it causes jitter.

m_particleMass Mass of a particle when interacting with rigid bodies and applying forces.

m_particleRadiusExpansion [world scale] Expands the collision detection radius of a particle. Can be used to find objects that are near but not colliding with the fluid particles. Setting this too high will reduce performance and degrade the accuracy of fluid-rigid contacts.

m_boundaryFriction [0.0, 1.0] Fraction of tangential velocity to remove per frame when a fluid particle collides with a rigid body. For example, using a m_boundaryFriction of 0.2 would cause a tangential velocity with magnitude 1 to be reduced to 0.8 after 1 frame. In practice, a value of e.g. 0.2 is already very high. Setting this too high will cause instability.

m_boundaryRestitution [0.0, 1.0] Generally set to 0 for fluids. It controls the fraction of normal velocity that is reflected when a fluid particle collides. Using a nonzero value will cause particles to bounce off rigid bodies when colliding. Note this introduces energy into the system, so it will make the fluid less stable.

m_boundaryErp Controls the amount of penetration removed per frame. To describe the collision resolution method, the force has 2 components. The first prevents the particle from further penetrating, by scaling as large as needed, and the second removes penetration. m_boundaryErp defines the second force.

3.3 Other Parameters

setCcdMotionThreshold The function `btCollisionObject::setCcdMotionThreshold()`, inherited by `btFluidSph`, can be used to prevent fluid particles from tunneling through rigid bodies. There is no feature to prevent rigid bodies from tunneling through fluids. In general, this should be set to `btFluidSphParametersLocal.m_particleRadius`.

4 Optimizations

Implementing a basic SPH algorithm is actually quite simple. Without any optimizations, it would simply take the form of a double nested loop. Much of the complexity comes from the grid algorithm.

4.1 Grid Algorithm

For spheres of the same radius, the most efficient algorithm is a uniform grid. There are several ways of implementing the uniform grid, some are presented here.

4.1.1 Static Noncolliding Uniform Grid (implemented by Bullet-Fluids)

It is a fixed-size sparse implicit grid, where only the contents of grid cells containing particles is stored.

The main advantage of this grid is that it has no collisions, and can be used in a fairly large world. The main disadvantage is that it requires a binary search to query a cell.

The information maintained by this algorithm includes:

- Cell size: A scalar. This is identical to the SPH interaction radius.
- value[]: An array containing *values* that identify grid cells. Each value is a 32-bit or 64-bit int. 10 or 21 bits are used for each dimension(x, y, z), so that worlds of 1024^3 and 2097152^3 cells can be used.
- index_range[]: An array containing the range of particle indices that corresponds to each grid cell. Each element contains 2 32-bit ints that store the first and last index of the particles in each grid cell.
- value_index[]: An array containing value-index pairs. Where the index identifies a particle and the *value* identifies the grid cell it is inside. There is one entry in this array for each particle.

value[] and index_range[] are parallel arrays, so value[i] identifies the grid cell with index i and contains index_range[i]. The length of the value[] and index_range[] array is equal to the number of grid cells with particles in them.

Grid Update Algorithm:

1. Quantize

First, a particle's position is converted into integer grid cell coordinates($\mathbf{R}^3 \rightarrow \mathbf{Z}^3$). This is achieved by dividing each coordinate of the particle's position by the grid cell width. Since a simple division would result in the cell (0,0,0) having a size of $2r$, we map the coordinate in range $[0, r)$ to 0, and $[-r, 0)$ to -1. Converts the continuous coordinates (x, y, z) to the integer coordinates $(x_{grid}, y_{grid}, z_{grid})$.

2. Assign value with hash function

Next, we combine these coordinates into a single value that can be used for sorting.

$$value = x_{grid} + y_{grid} * C + z_{grid} * C * C \quad (10)$$

where $C = 1024$ or 2097152 , defines the maximum extent of the grid in each dimension.

3. Create value-index pairs

We associate each particle with its value for sorting, creating an array of structs, each struct a value, index pair.

4. Sort by value-index pairs by value

On CPU, quick sort is used while radix sort is used on GPU. After sorting the particles are grouped by grid cells. The result of this step is an array of value-index pairs, where value_index[i] is the current position and value_index[i].index is the last position.

5. Rearrange particles

In order to avoid consuming a large amount of memory bandwidth during the sort, we only sort value-index pairs and then rearrange other particle attributes (position, velocity, user-pointer, etc.) after the sort.

6. Generate index ranges

Look through the particles to get the range of particle indices corresponding to each grid cell.

The result of this stage is a set of [value, index-ranges] pairs; the value is used to identify the grid cell, and the index-range indicates its contents.

Grid Query Algorithm:

1. Quantize

2. Assign value

In order to find neighbors, the first 2 steps is the same as the grid update.

3. Binary Search

Perform a binary search for the cell with value, once we have the index of the cell, it can be used to access the particle index range. Although a binary search may sound expensive, it only needs to be performed for each grid cell.

4.2 Grid Cell Size

The cell size $r * (3^3)$ is the optimal size for a CPU/GPU implementation. Compared to $2r * (2^3)$, it contains half the volume, so it is a closer fit to the actual neighbors.

On CPU, it also simplifies the symmetry optimization; since each particle in the same cell has the same 26 neighbor cells, it becomes possible to do processing on each grid cell instead of each particle. (It is actually possible to use the symmetry optimization with $2r$ cells, but it is much more complicated. The method requires making a copy of the grid and removing particles from the cell as the SPH density and SPH force is calculated. Additionally, since each $2r$ grid cell can still access 26 neighbors overall, it is still necessary to split the cells into 27 groups for multithreading, which reduces the amount of parallelism as each $2r$ cell has 8 times the volume of a r cell)

4.3 GPU Optimizations

On the GPU, we found that it is more efficient to store density and recompute the pressure instead of storing the inverted density and pressure. Using the ternary operator and moving the $(i \neq n)$ to the distance check also helps reduce divergence.

4.3.1 3-cell bar

If the hash function is defined as in the Bullet-Fluids grid(equation 10), it becomes possible to group 3 cells into a 3-cell bar, which reduces the amount of divergence. Instead of processing 27 cells, 9 3-cell bars are processed.

4.3.2 Effect of CPU optimizations on the GPU

Neighbor table Although using neighbor tables improves performance on CPU, it decreases overall performance and increases memory consumption on the GPU. To be exact, it reduces the performance of the density computation stage and improves the performance of the force calculations stage. There may be a performance benefit for methods that iterate over all particles several times, however.

0							7
8			11	12	13		15
16			19	20	21		23
24			27	28	29		31
32							39
40							47
48							55
56							63

19	20	21
[22, 32]	[33, 41]	[42, 54]

Figure 5: 2D example of 3-cell bar optimization. In this case, the constant C in equation 10 is equal to 8. Instead of using a for loop with 9 iterations(27 in 3D), it is possible to use a for loop with 3 iterations(9 in 3D), with each iteration processing 3 cells. In this example, the cells with values 19, 20, and 21 have particles with indices [22, 32], [33, 41], and [42, 54], respectively. Instead of running 3 iterations for each of these particle index ranges, we can perform a single iteration through the range [22, 54].

Symmetry Implementing the symmetry optimization as on CPU was found to result in a $\sim 50x$ performance decrease. This is likely because:

- Processing by grid cell requires more registers.
- Higher memory latency on GPU means that the symmetric optimization is less effective.
- There are less grid cells than particles, which reduces the max number of active threads. Additionally, grouping the cells into 27 groups further reduces the occupancy. (A GPU typically needs 1000–10000 or more threads for efficient processing. On average, a grid cell with size r has 4 particles and 27 groups are needed to avoid thread collisions, so the overall number of threads that can be used is reduced by an estimate of $4 * 27 = 108$.

4.4 CPU Optimizations

4.4.1 Symmetry and Multithreading

In order to allow multithreading without using atomics or other synchronization, the grid cells are divided into 27 groups (9 groups in 2D) so that processing all the cells in a single group will not cause thread collisions.

Since the pressure and viscosity forces are antisymmetric, that is, $\mathbf{f}_{ij} = -\mathbf{f}_{ji}$, it is only necessary to compute forces once for each pair of particles. The density contribution between particles is also symmetric, so the same optimization can be applied during the density computation stage. This additionally reduces the number of neighbor cells that needs to be accessed by each cell from 26 to 13 cells in 3D and 8 to 4 cells in 2D.

1	2	3	1	2	3
4	5	6	4	5	6
7	8	9	7	8	9
1	2	3	1	2	3
4	5	6	4	5	6
7	8	9	7	8	9

1	2	3	1	2	3
4	5	6	4	5	6
7	8	9	7	8	9
1	2	3	1	2	3
4	5	6	4	5	6
7	8	9	7	8	9

Figure 6: Multithreading optimization. A 2D example is presented for 2 iterations. The cells are grouped into clusters of 9 cells, and due to symmetry each cell only needs to access 4 neighbor cells during the computation. During the first iteration, cells marked 1 can be processed by multiple threads. During the second iteration, cells marked 2 are processed, and so on. Dark blue indicates the current cell, while light blue indicates neighbor cells that are accessed and written to.

4.4.2 Neighbor Tables

The index and distance of neighbor particles are stored during the SPH density calculation to avoid recomputing during the SPH force calculation. Due to symmetry, each entry represents a pair. This also allows tuning the quality of the simulation by reducing the max number of neighbors.

5 Other Internal Details

5.1 Update Loop

1. Update grid
2. Calculate SPH sum (density and pressure)
3. Calculate SPH pressure and viscosity force
4. Integrate velocity (apply SPH forces)
5. Handle collision detection and response with rigid bodies (correct velocities to avoid penetrating into rigids)
6. Integrate position

5.2 Rigid Body Interaction

Fluid-Rigid interaction is implemented by treating particles as rigid body spheres. CCD is implemented for moving fluid against rigid by casting a ray along the path of a fluid particle's motion. If the ray intersects a rigid body, then the particle is moved to intersect slightly with the rigid body. The collision response step then prevents the particle from tunneling by applying an impulse that is as large as needed.

A drawback of this method is that fluid particles can stack while colliding with rigid bodies. That is, since the fluid is not completely incompressible, many fluid particles can occupy the same space and, as a result, cause a colliding rigid body to experience very strong collision forces. This issue is reduced somewhat if a less compressible, but more expensive, fluid simulation method is used.

Another limitation is that the SPH density is inaccurate at boundaries, as SPH assumes that each particle has a full neighborhood. This causes fluids to cluster at boundaries in order to make up for the reduction in density.

Dynamic rigid bodies, especially triangle meshes, are able to tunnel through the fluid if moving too fast. Currently, there is no solution for this issue aside from reducing the time step.

5.3 Other notes

Clustering When implementing a new SPH solver, a common issue is that the particles only seem to attract each other and form into clumps. This may be caused by using the incorrect sign for the pressure force (or a negative stiffness).

Split masses The mass of a particle is different when calculating the SPH force and when interacting with rigid bodies. This allows the user to separately tune the interaction with rigid bodies. Another motivation is that large mass ratios between rigid bodies can cause the simulation to become unstable.

Local and Global parameters split In order to efficiently perform fluid-fluid interaction, it is necessary to use a common smoothing radius and simulation scale. Otherwise, the grids would be out of alignment and a more complex approach that is also slower would be needed.

Denormalized floating point The variables used by SPH can become very close to 0, causing the floating point values to become denormalized. This has the potential to cause extreme degradations in performance, although the author has not encountered such issues even with 32-bit floats.

References

- [1] Erwin Coumans et al. Bullet physics library (version 2.81). <http://bulletphysics.org>, 2012. Accessed: 30 December 2013.
- [2] Rama Hoetzlein. Fluids v.2 - sph fluid simulator for cpu and gpu. <http://www.rchoetzlein.com/eng/graphics/fluids.htm>, 2008. Accessed: 30 December 2013.
- [3] Micky Kelager. Lagrangian fluid dynamics using smoothed particle hydrodynamics. Master’s thesis, University of Copenhagen, Department of Computer Science, January 2006.
- [4] Matthias Müller, David Charypar, and Markus Gross. Particle-based fluid simulation for interactive applications. In *Proceedings of the 2003 ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, SCA ’03, pages 154–159. Eurographics Association, 2003.